

---

## CSE 150A/250A - Homework 4

**Out:** *Tues Oct 22*

**Due:** *Mon Oct 28* (by 11:59 PM, Pacific Time, via Gradescope)

**Grace period:** 24 hours

### General Guidelines:

You should submit your homework assignments via Gradescope. Typesetting  $\text{\LaTeX}$  is preferred, but neatly handwritten solutions are also accepted. Upload a PDF of your answers to Homework 4 on Gradescope, and carefully select the pages corresponding to each question using the submission interface. If you are submitting handwritten answers, please scan them and create a PDF for upload. Here is a primer on submitting PDF homework via Gradescope:

<https://tinyurl.com/gradescope-guide>

If you have not done this before, please allow some extra time to familiarize yourself with this process.

**Late Day Policy:** There will be no penalty for turning in any of the homework assignments up to 24 hours late. However, we cannot guarantee that late assignments will be graded in a timely fashion. Beyond the 24-hour grace period, we will not accept late homework.


**Collaboration:** We strongly encourage collaboration (**but NOT copying**) on the homework assignments. You may talk to anyone in the course about how to solve the problems, and you may even compare your solutions. However, you must write up your solutions yourself, and you may not copy them.

---

---


## 4.1 Maximum likelihood estimation in belief networks (8 pts)

Consider the two DAGs shown below,  $G_1$  and  $G_2$ , over the same nodes  $\{X_1, X_2, \dots, X_n\}$ , that differ only in the direction of their edges.



Suppose that we have a (fully observed) data set  $\{x_1^{(t)}, x_2^{(t)}, \dots, x_n^{(t)}\}_{t=1}^T$  in which each example provides a complete instantiation of the nodes in these DAGs. Let  $\text{COUNT}_i(x)$  denote the number of examples in which  $X_i = x$ , and let  $\text{COUNT}_i(x, x')$  denote the number of examples in which  $X_i = x$  and  $X_{i+1} = x'$ .

- Express the maximum likelihood estimates for the CPTs in  $G_1$  in terms of these counts. (2 pts)
- Express the maximum likelihood estimates for the CPTs in  $G_2$  in terms of these counts. (2 pts)
- Using your answers from parts (a) and (b), show that the maximum likelihood CPTs for  $G_1$  and  $G_2$  from this data set give rise to the same joint distribution over the nodes  $\{X_1, X_2, \dots, X_n\}$ . (2 pts)
- Suppose that some but not all of the edges in these DAGs were reversed, as in the graph  $G_3$  shown below. Would the maximum likelihood CPTs for  $G_3$  also give rise to the same joint distribution? (Hint: does  $G_3$  imply all the same statements of conditional independence as  $G_1$  and  $G_2$ ? Look in particular at nodes  $X_3$  and  $X_{n-2}$ .) (2 pts)



---

## 4.2 Statistical language modeling (15 pts)

In this problem, you will explore some simple statistical models of English text. Download and examine the data files on Canvas for this assignment. (Start with the `hw4_readme.txt` file.) These files contain unigram and bigram counts for 500 frequently occurring tokens in English text. These tokens include actual words as well as punctuation symbols and other textual markers. In addition, an “unknown” token is used to represent all words that occur outside this basic vocabulary. For this problem, as usual, you may program in the language of your choice.

- (a) Compute the maximum likelihood estimate of the unigram distribution  $P_u(w)$  over words  $w$ . Print out a table of all the tokens (i.e., words) that start with the letter “M”, along with their numerical unigram *probabilities* (not counts). (You do not need to print out the unigram probabilities for all 500 tokens.) (1 pt)
- (b) Compute the maximum likelihood estimate of the bigram distribution  $P_b(w'|w)$ . Print out a table of the ten most likely words to follow the word “THE”, along with their numerical bigram probabilities. (1 pt)
- (c) Consider the sentence “**The stock market fell by one hundred points last week.**” Ignoring punctuation, compute and compare the log-likelihoods of this sentence under the unigram and bigram models:

$$\begin{aligned}\mathcal{L}_u &= \log \left[ P_u(\text{the}) P_u(\text{stock}) P_u(\text{market}) \dots P_u(\text{points}) P_u(\text{last}) P_u(\text{week}) \right] \\ \mathcal{L}_b &= \log \left[ P_b(\text{the}|\langle s \rangle) P_b(\text{stock}|\text{the}) P_b(\text{market}|\text{stock}) \dots P_b(\text{last}|\text{points}) P_b(\text{week}|\text{last}) \right]\end{aligned}$$

In the equation for the bigram log-likelihood, the token  $\langle s \rangle$  is used to mark the beginning of a sentence. Which model yields the highest log-likelihood? (3 pts)

- (d) Consider the sentence “**The sixteen officials sold fire insurance.**” Ignoring punctuation, compute and compare the log-likelihoods of this sentence under the unigram and bigram models:

$$\begin{aligned}\mathcal{L}_u &= \log \left[ P_u(\text{the}) P_u(\text{sixteen}) P_u(\text{officials}) \dots P_u(\text{sold}) P_u(\text{fire}) P_u(\text{insurance}) \right] \\ \mathcal{L}_b &= \log \left[ P_b(\text{the}|\langle s \rangle) P_b(\text{sixteen}|\text{the}) P_b(\text{officials}|\text{sixteen}) \dots P_b(\text{fire}|\text{sold}) P_b(\text{insurance}|\text{fire}) \right]\end{aligned}$$

Which pairs of adjacent words in this sentence are not observed in the training corpus? What effect does this have on the log-likelihood from the bigram model? (3 pts)

- (e) Consider the so-called *mixture* model that predicts words from a weighted interpolation of the unigram and bigram models:

$$P_m(w'|w) = \lambda P_u(w') + (1 - \lambda) P_b(w'|w),$$

where  $\lambda \in [0, 1]$  determines how much weight is attached to each prediction. Under this mixture model, the log-likelihood of the sentence from part (d) is given by:

$$\mathcal{L}_m = \log \left[ P_m(\text{the}|\langle s \rangle) P_m(\text{sixteen}|\text{the}) P_m(\text{officials}|\text{sixteen}) \dots P_m(\text{fire}|\text{sold}) P_m(\text{insurance}|\text{fire}) \right].$$

Compute and plot the value of this log-likelihood  $\mathcal{L}_m$  as a function of the parameter  $\lambda \in [0, 1]$ . From your results, deduce the optimal value of  $\lambda$  to two significant digits. (3 pts)

- (f) Submit a printout of your source code for the previous parts of this problem. (4 pts)
-



(c) **Likelihoods**

Consider again the training sequence  $\mathcal{S}$ , as well as three test sequences  $\mathcal{T}_1$ ,  $\mathcal{T}_2$ , and  $\mathcal{T}_3$  of the same length, shown below. Note that  $\mathcal{T}_2$  and  $\mathcal{T}_3$  contain bigrams (underlined) that are not in the training sequence  $\mathcal{S}$ .

$\mathcal{S}$  = "a a b b b b c c d d a a d d c c"  
 $\mathcal{T}_1$  = "a d a d a d a d a d a d a d"  
 $\mathcal{T}_2$  = "a a a d d d d c c c c b b b b"  
 $\mathcal{T}_3$  = "b d b d b d b d b d b d b d"

Consider the probabilities of these sequences under the unigram and bigram models from parts (a) and (b) of this problem (i.e., the models that you estimated from the training sequence  $\mathcal{S}$ ). For each of the following, indicate whether the probability on the left is equal (=), greater (>), or less (<) than the probability on the right.

Note: you can (and should) answer these questions without explicitly computing the numerical values of the expressions on the left and right hand sides.

$P_U(\mathcal{S})$    $P_U(\mathcal{T}_1)$

$P_U(\mathcal{S})$    $P_U(\mathcal{T}_2)$

$P_U(\mathcal{S})$    $P_U(\mathcal{T}_3)$

$P_B(\mathcal{T}_1)$    $P_B(\mathcal{S})$

$P_B(\mathcal{T}_2)$    $P_B(\mathcal{S})$

$P_B(\mathcal{T}_3)$    $P_B(\mathcal{T}_2)$

$P_U(\mathcal{S})$    $P_B(\mathcal{S})$

$P_U(\mathcal{T}_1)$    $P_B(\mathcal{T}_1)$

$P_U(\mathcal{T}_2)$    $P_B(\mathcal{T}_2)$

$P_U(\mathcal{T}_3)$    $P_B(\mathcal{T}_3)$

(d) **Likelihoods**

Consider the model obtained by linear interpolation (or mixing) of the unigram and bigram models estimated in part (a) of this problem:

$$P_M(\tau'|\tau) = (1 - \lambda)P_1(\tau') + \lambda P_2(\tau'|\tau),$$

with mixing coefficient  $\lambda \in [0, 1]$ . For a sequence of tokens of length  $L$ , the mixture model computes the log-likelihood as:

$$\mathcal{L} = \log P_1(\tau_1) + \sum_{\ell=2}^L \log P_M(\tau_\ell|\tau_{\ell-1}).$$

Naturally, this value varies as a function of the coefficient  $\lambda$ . For  $\lambda$  near zero, it is close to the log-likelihood of the unigram model; for  $\lambda$  near one, it is close to that of the bigram model. This last part of this problem asks you to consider, for each of the sequences below, the *qualitative* behavior of the mixture model's log-likelihood as a function of  $\lambda \in [0, 1]$ . (For instance, is this function constant, or if not, where do its maximum and minimum occur?)

The plots below illustrate four possible behaviors of the mixture model's log-likelihood as a function of  $\lambda \in [0, 1]$ . For each sequence below, indicate the one plot (either A, B, C, or D) that sketches the correct qualitative behavior. (Note that these graphs are not exactly what they would look like if we were to graph the function; we are only asking about the general behavior of the function (e.g. increasing as  $\lambda$  goes to 1.))

$\mathcal{S} =$  "a a b b b b b c c d d a a d d c c"

☐

$\mathcal{T}_1 =$  "a d a d a d a d a d a d a d a d"

☐

$\mathcal{T}_2 =$  "a a a a d d d d c c c c c b b b b"

☐

$\mathcal{T}_3 =$  "b d b d b d b d b d b d b d b d"

☐

